JobMatchHire — ECCE Evaluation Report

Explainable Contextual Competency Engine with Automated Counterfactual Gap Remediation

1. Evaluation Overview & Identification

Job Position:	Staff Cloud Systems Architect	**Evaluation Date (IST):**	2026-10-25T17:46:06Z
Candidate Name:	Emily Watson	**Candidate ID:**	CAN-725D4D
Target Job ID:	JOB-365	**Evaluation Status:**	**Evaluation Completed**

2. Overall Match Score & Deterministic Component Breakdown

Component	Weight / Max	Score Achieved	Metric / Formula
Required Skills	50.0 pts	**12.50 pts**	Ratio of verified required skills to total required
Professional Experience	25.0 pts	**25.00 pts**	Clamped ratio of verified years vs minimum requirement
Preferred Skills	15.0 pts	**0.00 pts**	Ratio of verified preferred skills to total preferred
Text Similarity (TF-IDF Cosine)	10.0 pts	**0.31 pts**	Cosine similarity score: 0.031
TOTAL MATCH SCORE	**100.0 pts**	**37.81 / 100**	**Deterministic Bounded Sum [0.0 - 100.0]**

3 & 4. Competency Assessment & Evidence Classification

*ECCE classifies evidence strictly based on extracted resume passages. **Evidence Not Found** must not be interpreted as proof that the candidate lacks the skill. **Potential Inconsistency** is flagged for human review, not treated as proof of dishonesty.*

Skill	Type	Classification	Detected Actions & Project Context	Assessment Notes
Python	Required	**Evidence Found**	**Actions:** None identified **Context:** General mention	Verified supporting context with active project application.
AWS	Required	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.
Kubernetes	Required	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.
Docker	Required	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.

PostgreSQL	Preferred	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.
Terraform	Preferred	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.
Redis	Preferred	**Evidence Not Found**	**Actions:** None identified **Context:** General mention	Skill is not explicitly referenced in the extracted resume text.

5. Supporting Verbatim Resume Passages

Verbatim excerpts from candidate's submitted resume documenting contextual competency claims without fabrication:

Skill: Python (Required)

"Languages: Python, JavaScript, HTML, CSS, SQL"

"• Assisted in creating Python Flask web services with SQLite databases"

6. Scoring Explanation & Role Gap Analysis

Evaluator Summary: Verified 1/4 required competencies with contextual evidence; missing: AWS, Docker, Kubernetes. Experience: Exceeds Requirement (7.0y vs 4.0y req).

7. Counterfactual Score Simulations (Hypothetical Analysis)

NOTICE: The following simulations represent hypothetical mathematical sensitivity models under the deterministic scoring function. They illustrate how specific changes in verified qualifications or resume parsing clarification would affect the matching score. These are hypothetical models, not actual results or awarded points. Candidates must never add skills they do not genuinely possess.

Scenario	Type	Baseline	Hypothetical	Impact / Explanation
Competency Acquisition: AWS	Skill Acquisition	37.81%	**50.31%** (+12.50%)	Candidate undertakes formal training or practical project work to develop verifiable competency in AWS. *Hypothetical Simulation - Not Actual Result*
Parsing/Format Clarification: AWS	Resume-Parsing Correction	37.81%	**50.31%** (+12.50%)	If candidate already possesses verifiable hands-on AWS experience that was omitted or formatted in non-extractable text, clarifying the resume with active project bullets adjusts the score. *Hypothetical Simulation - Not Actual Result*
Preferred Skill Acquisition: PostgreSQL	Preferred Skill Acquisition	37.81%	**42.81%** (+5.00%)	Candidate demonstrates complementary competence in preferred skill PostgreSQL. *Hypothetical Simulation - Not Actual Result*

8. Candidate Skill-Gap Guidance & Practical Next Steps

[PRIVATE TO CANDIDATE] The following development recommendations are provided exclusively to assist the candidate in closing identified capability gaps. These recommendations are kept strictly private and are not disclosed to recruiters.

Skill Gap	Relevance to Job Role	Recommended Practical Project	Authoritative Resource
AWS (Required)	Industry-standard cloud platform for hosting scalable infrastructure, managed storage, and serverless compute.	Deploy a containerized application to AWS ECS or EKS with RDS database integration, S3 storage, and CloudWatch monitoring.	AWS Well-Architected Framework & AWS Certified Solutions Architect Learning Path
Docker (Required)	Essential for containerization, environment reproducibility, and microservice orchestration across environments.	Containerize a multi-tier web application using multi-stage Docker builds, environment variables, and Docker Compose networking.	Docker Official Documentation (docs.docker.com) & Best Practices for Containerization
Kubernetes (Required)	Industry standard for automated deployment, scaling, and operational management of containerized workloads.	Configure a local Kubernetes cluster (Minikube/k3s) with Deployments, Services, ConfigMaps, Ingress, and Horizontal Pod Autoscaling.	Kubernetes Official Interactive Tutorial (kubernetes.io/docs/tutorials)
PostgreSQL (Preferred)	Key technical capability specified in the job requisitions for preferred qualifications.	Build a practical proof-of-concept project demonstrating production usage of PostgreSQL.	Official documentation and industry-standard tutorials for PostgreSQL.
Terraform (Preferred)	Key technical capability specified in the job requisitions for preferred qualifications.	Build a practical proof-of-concept project demonstrating production usage of Terraform.	Official documentation and industry-standard tutorials for Terraform.
Redis (Preferred)	In-memory data store essential for sub-millisecond caching, rate-limiting, and distributed session management.	Implement a distributed cache and sliding-window rate limiter in front of a relational database with TTL invalidation policies.	Redis Official Documentation (redis.io/docs) & Caching Patterns Guide

---

Data Distinction Standard: Verified inputs include job descriptions, candidate submitted files, and deterministic score formulas. Resume claims reflect candidate-supplied text. Contextual models classify evidence. Counterfactual simulations are strictly hypothetical.